




Fixing the E3xx DMA Bottleneck

Implementing a High-Rate Heterogeneous FPGA/Processor Transport

EJ Kreinar
Lorin Metzger
Dan CaJacob

Topics

- ▶ Problem Introduction & Motivation
- ▶ Direct Memory Access (DMA) Background
- ▶ Technical Solution
 - FPGA
 - Kernel Driver
 - User-Space Software
- ▶ Results & Tradeoffs

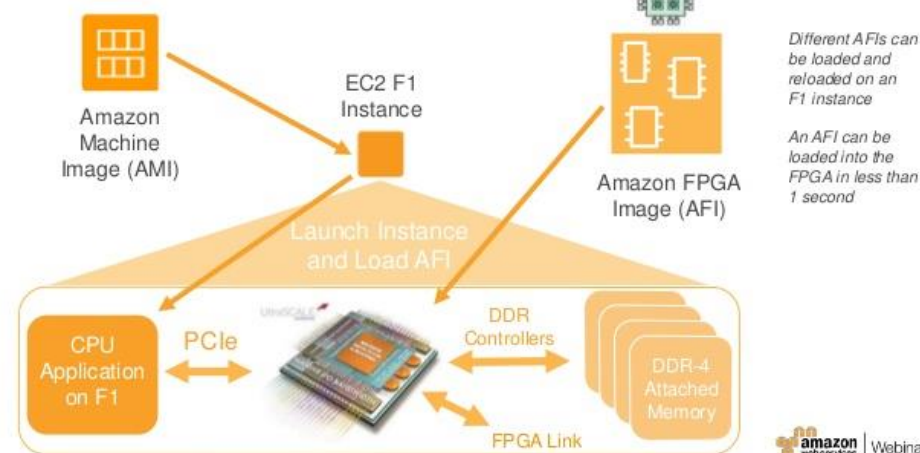


- ▶ Heterogeneous computing requires a data transfer
 - Data must be moved from one computing device to another
- ▶ This data transfer is often a bottleneck in system throughput and latency




NVIDIA GPU

FPGA Acceleration Using F1



AWS F1 CPU + FPGA



"Aller" Artix 7 FPGA with M.2 Interface

Takeaway Point

- ▶ Self-contained, low-SWAP Software Defined Radios...



Ettus E310



Ettus E320



Epiq Sidekick Z2



Lime SDR

- ▶ Often using a Zynq "System on Chip": FPGA + ARM Processor

We'll focus on the Xilinx Zynq platform

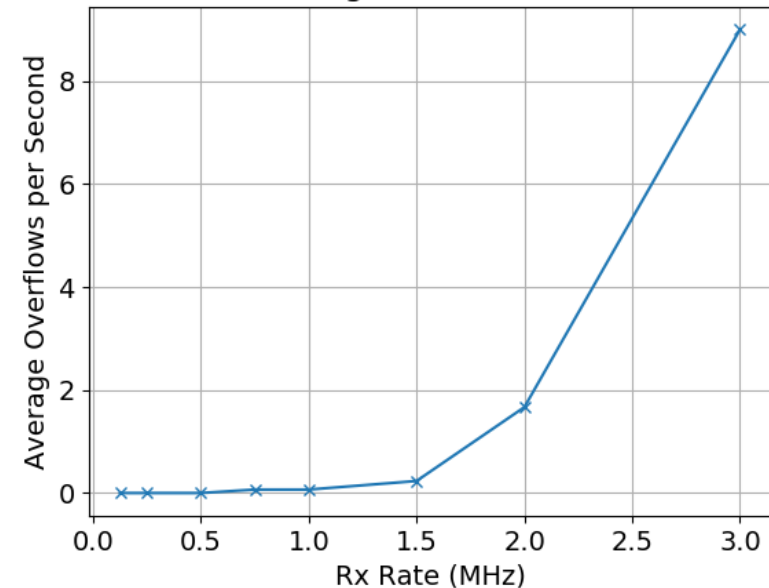
▶ Example Device: Ettus E310

- (No PS Load) Max continuous receive rate without overflows: **10 MHz**
- (Writing to File) Max continuous receive rate without overflow: **~1 MHz**

▶ File Write Test:

- 1-channel receive only. Increment rx sample rate
- Collect continuously for 30 seconds
- Write to disk in natural "wire format", ie interleaved short 16-bit ints

Overflow Rate during E310 Rx to File (30 Seconds)

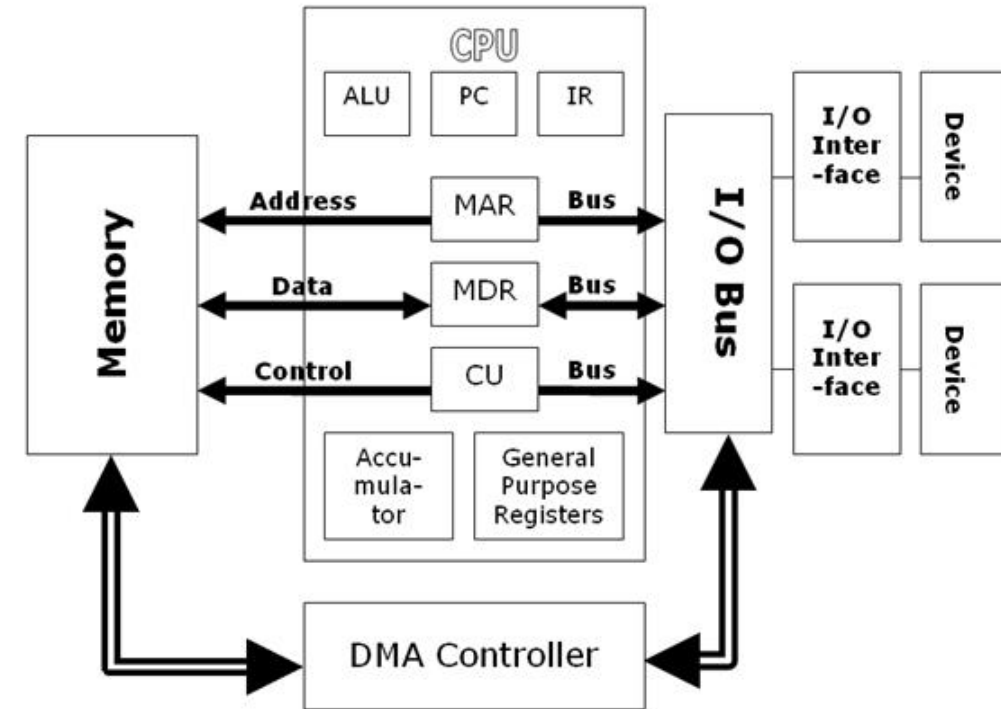
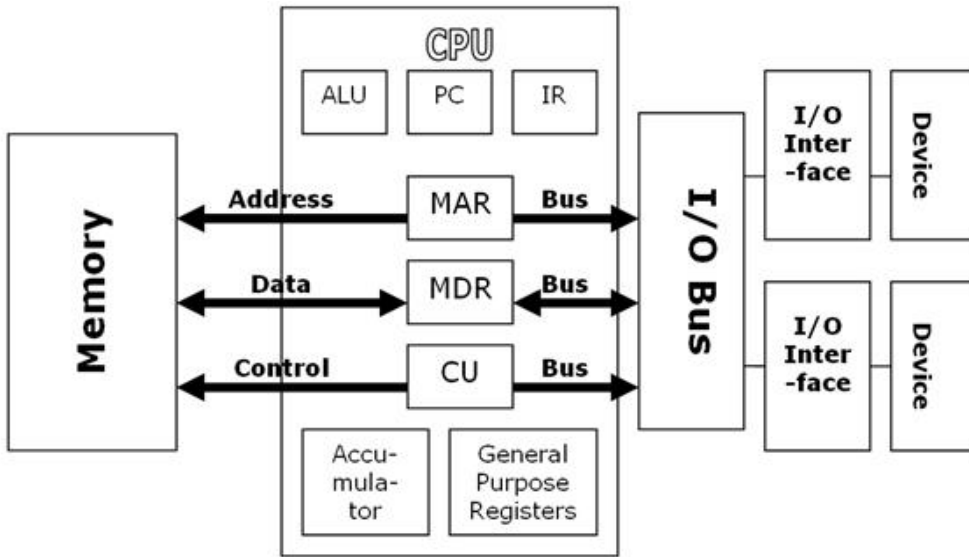


Some devices may NOT achieve theoretical limits



Direct Memory Access (DMA) Background

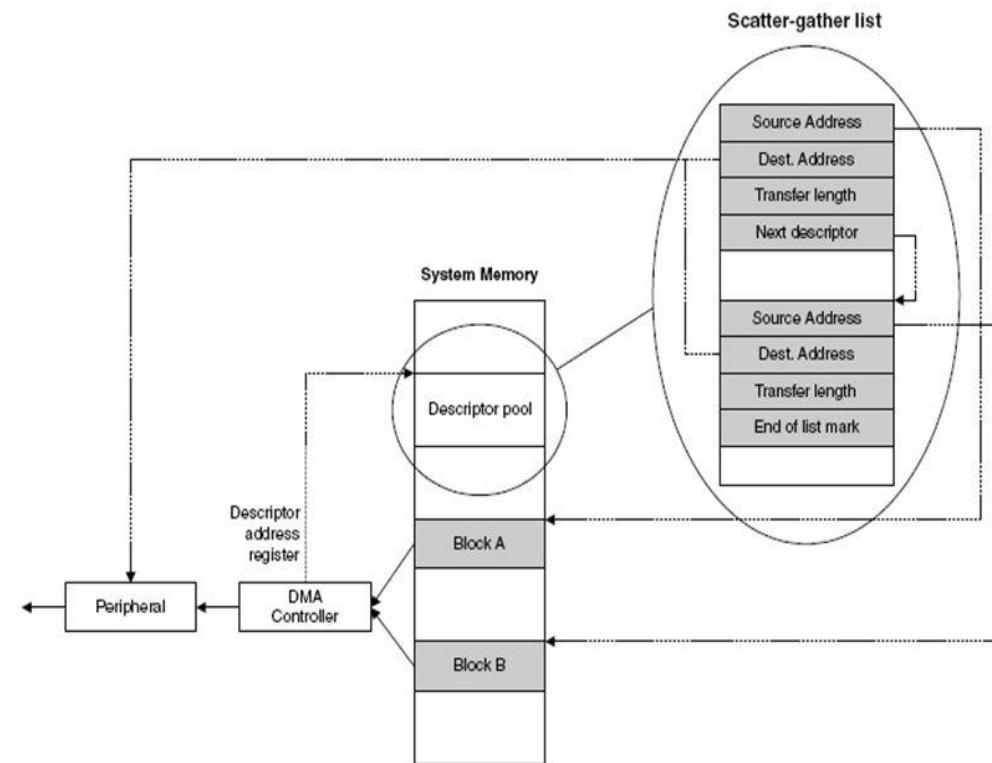
- ▶ Direct Memory Access is a feature that allows memory copy operations to happen *without* processor involvement



DMA generally requires use of a purpose-designed hardware architecture to perform the optimized copy

- ▶ Scatter-gather allows the transfer of data to and from multiple memory areas in a single DMA transaction
- ▶ It is equivalent to chaining together multiple simple DMA requests
 - The motivation is to off-load multiple input/output interrupt and data copy tasks from the CPU
 - Allows for more efficient memory use by kernel by with less potential memory fragmentation vs large continuous buffers.

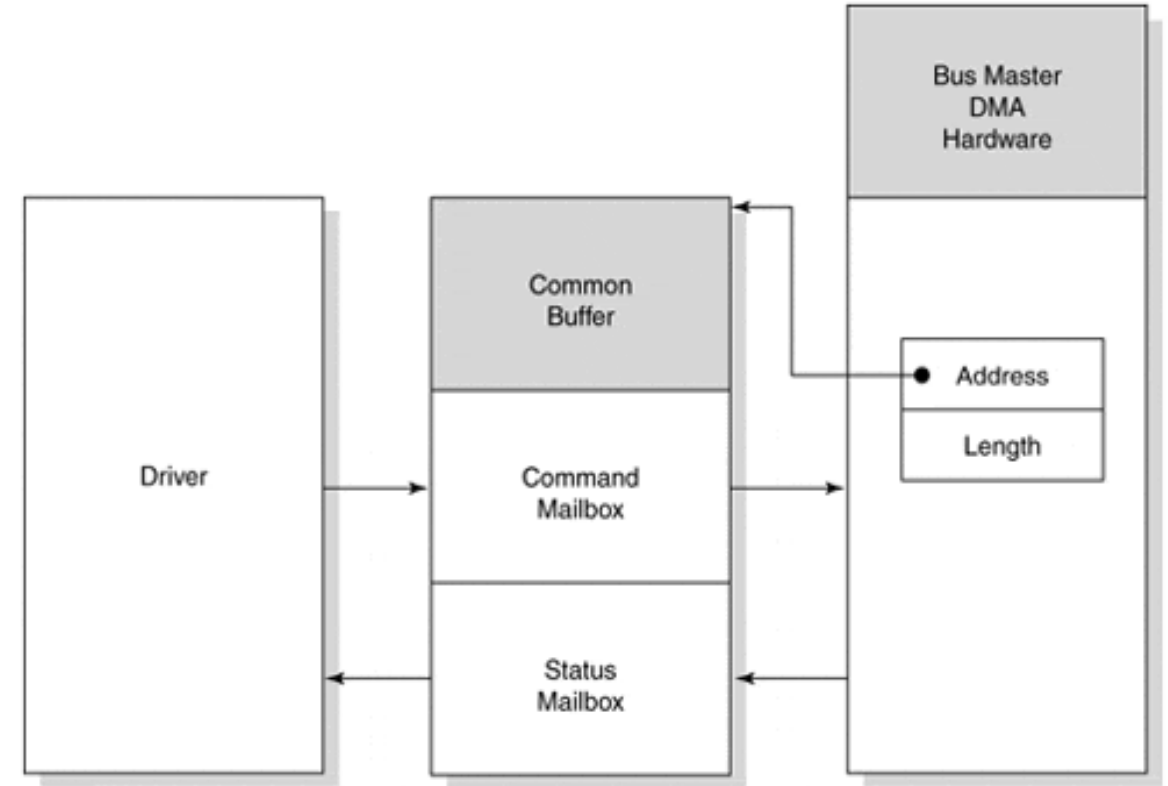
Scatter Gather DMA Mode



Source: Symbian OS Internals/13. Peripheral Support

FPGA controls buffer swapping with minimal PS interaction

- ▶ Direct Register Mode (non-scatter gather) provides a configuration for doing simple DMA transfers
- ▶ Transfers are initiated by accessing the DMACR, the Source or Destination Address, and the Length registers
- ▶ When the transfer is completed, an IRQ register asserts for the associated channel and if enabled generates an interrupt output signal



PS required to service IRQs and manipulate DMA buffer

- ▶ FPGA fabric is allocated on the system's AXI4 crossbar
- ▶ FPGA IP interacts with the DDR Memory Controller through the memory interconnects

Table 22-2: Theoretical Bandwidth of PS-PL and PS Memory Interfaces

Interface	Type	Bus Width (bits)	IF Clock (MHz)	Read Bandwidth (MB/s)	Write Bandwidth (MB/s)	R+W Bandwidth (MB/s)	Number of Interfaces	Total Bandwidth (MB/s)
General Purpose AXI	PS Slave	32	150	600	600	1,200	2	2,400
General Purpose AXI	PS Master	32	150	600	600	1,200	2	2,400
High Performance (AFI) AXI_HP	PS Slave	64	150	1,200	1,200	2,400	4	9,600
AXI_ACP	PS Slave	64	150	1,200	1,200	2,400	1	2,400
DDR	External Memory	32	1,066	4,264	4,264	4,264	1	4,264
OCM	Internal Memory	64	222	1,779	1,779	3,557	1	3,557

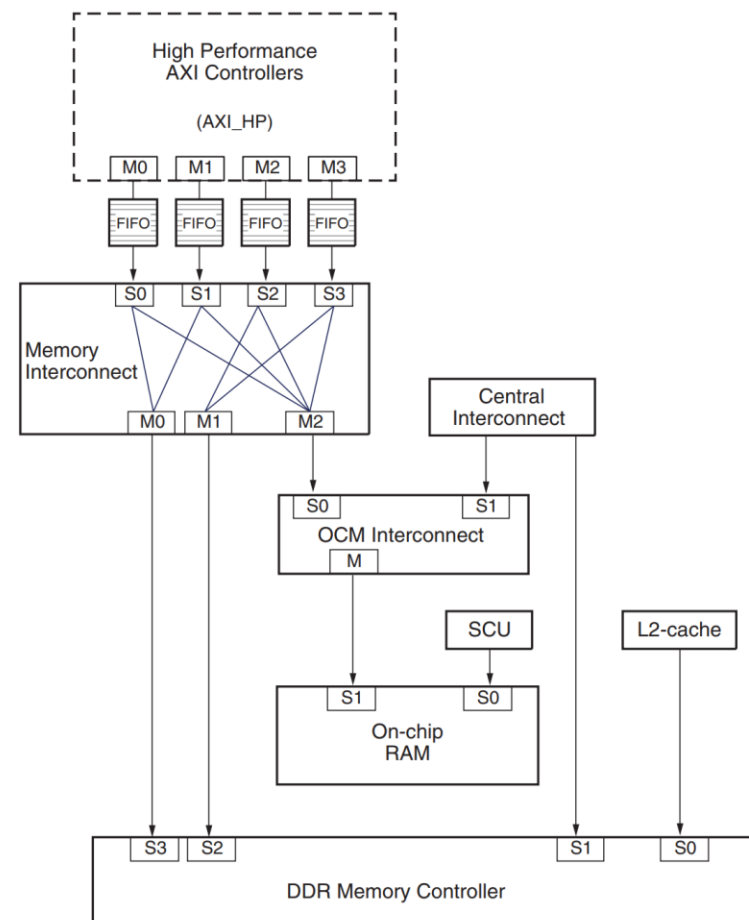
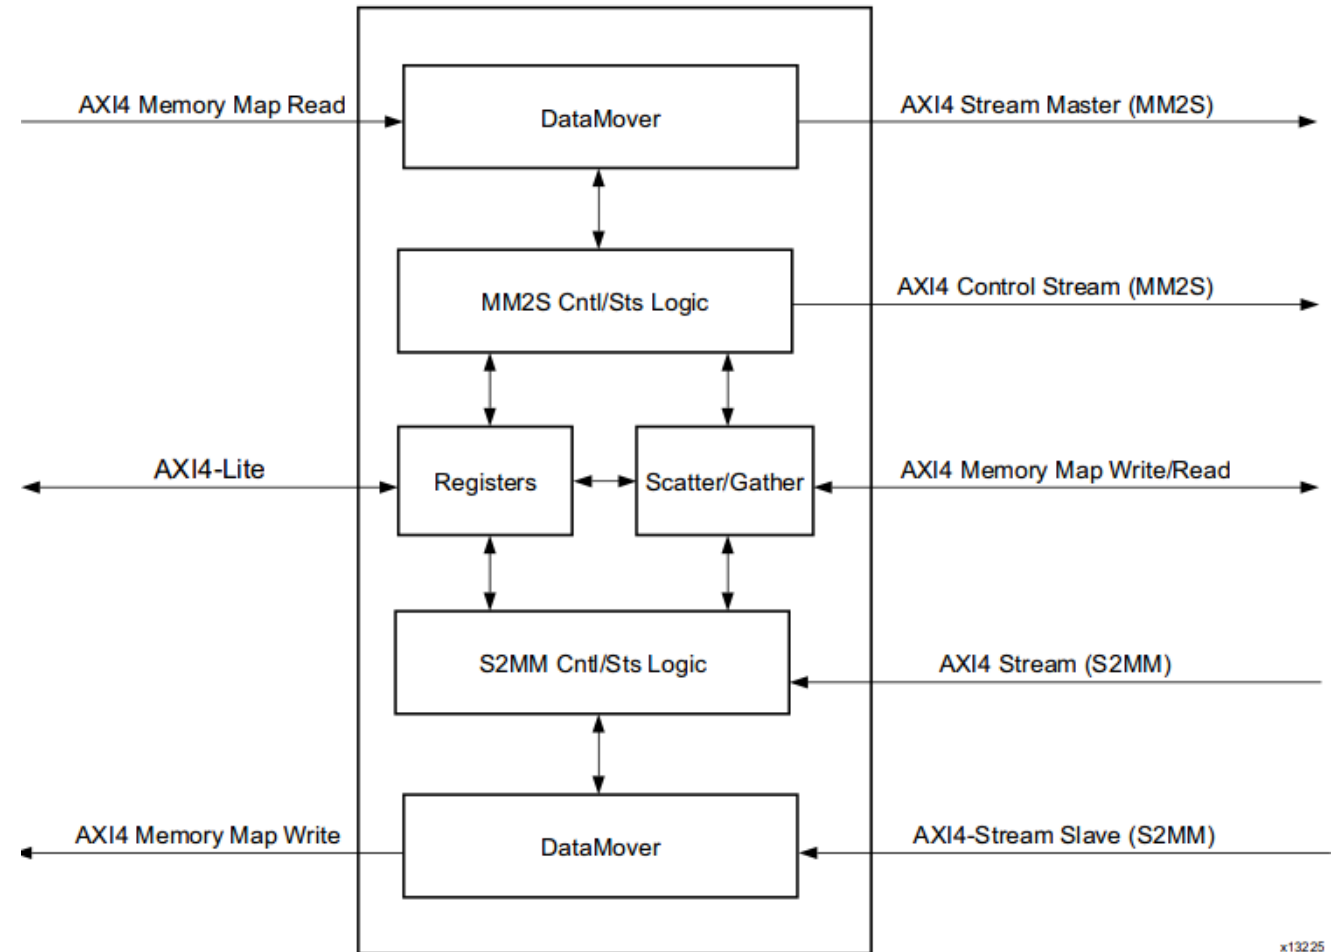


Figure 5-4: High Performance (AXI_HP) Connectivity

Max theoretical bandwidth of 600 MB/s - 1200 MB/s

- ▶ AXI DMA distinguishes two channels
 - MM2S (memory-mapped to stream) transports data from DDR memory to FPGA
 - S2MM (stream to memory-mapped) transports arbitrary data stream to DDR memory
- ▶ Uses AXI-4-Lite interface for configuring DMA control registers.



Polling driver periodically interacts with DMA registers

▶ Benefits:

- Simple to implement and test!

▶ Drawbacks:

- Must run PS in a tight loop to respond quickly to DMA

Interrupt driver interacts with DMA on demand based on IRQ signals

▶ Benefits:

- Allows the PS to react immediately to interrupt request signals
- Software does not load the PS if DMA is not used

▶ Drawbacks:

- Requires kernel module development. Userspace programs do not directly handle interrupts in a Linux OS

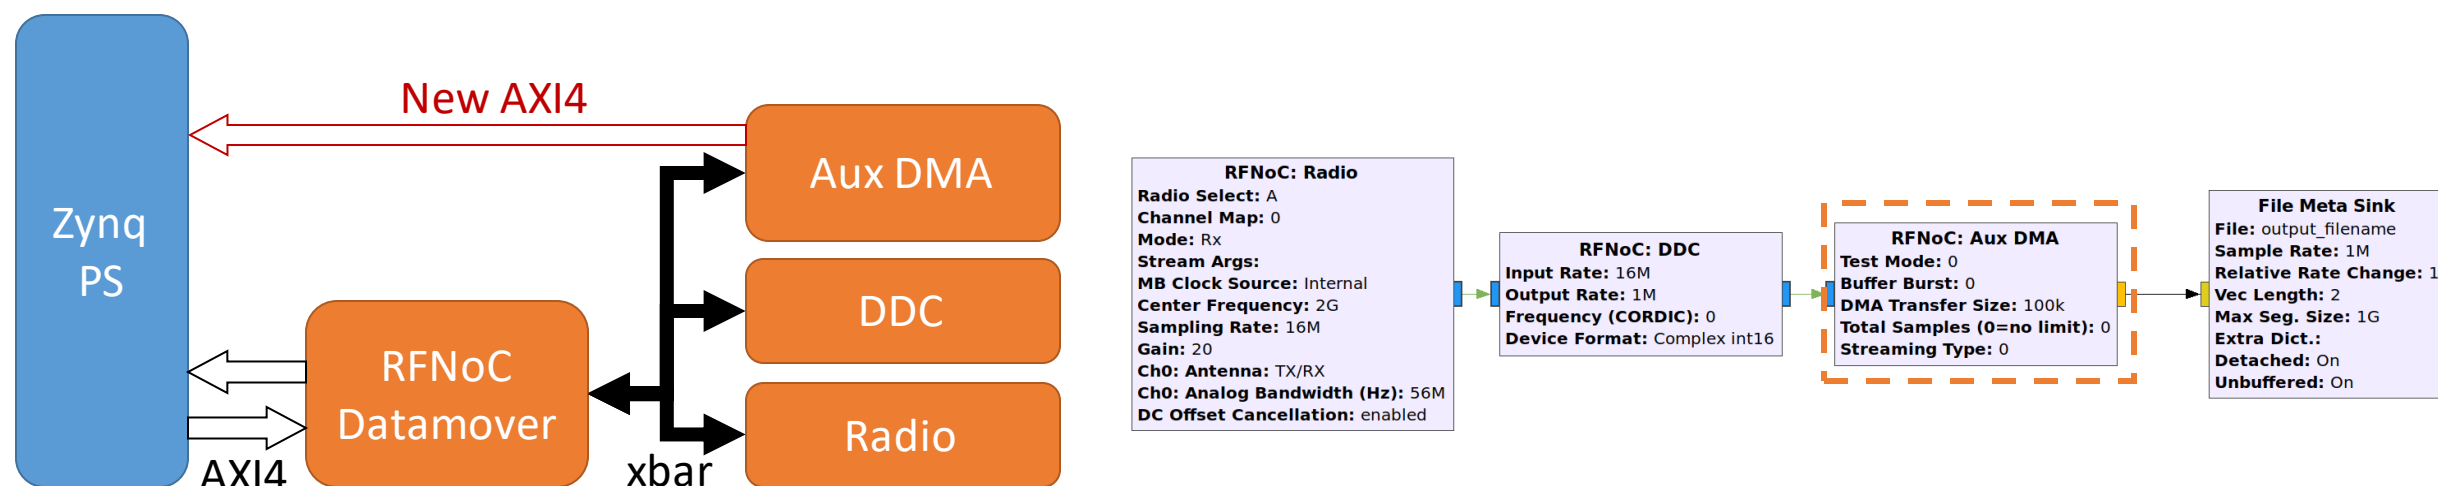
Polling may be helpful to prototype, but mostly all DMA software ultimately requires a kernel driver with interrupt handling



Technical Solution:
FPGA, Kernel Driver, User-Space Software

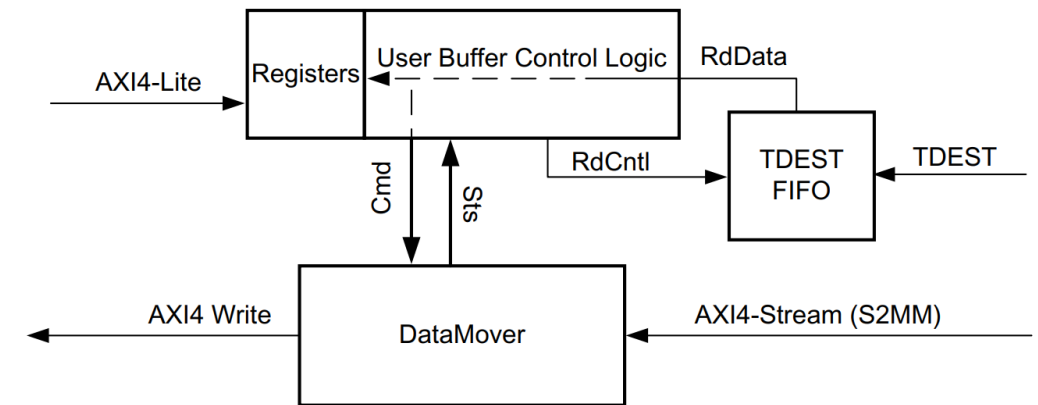

Major Surgery:

- ▶ Aux DMA: `noc_block_auxdma` connects directly to the E310 processor via AXI4 DMA
- ▶ New kernel driver services DMA interrupts
- ▶ Reuse typical RFNoC register reads/writes and graph infrastructure
- ▶ Aux DMA overrides "general_work" function of `rfnoc_block_impl` (in `gr-ettus`)



"Aux DMA" sends data directly into shared processor DRAM using large packet sizes

- ▶ Xilinx AXI Datamover: Core data transfer functionality
 - Writes "Serial Data" to "Memory-Mapped" data (and vice versa) using the FPGA as the AXI-Master
- ▶ Xilinx AXI DMA
 - Wraps Datamover and adds control registers for Scatter Gather or Direct Register Modes
- ▶ RFNoC Integration: **Use Xilinx AXI DMA**
 - Create a new streaming "endpoint"
 - Transfer from "start time" to "end of burst"



X13067

Figure 3-2: Typical Application of S2MM DataMover

▶ Xilinx AXI DMA Linux Driver

- Implemented as a Linux DMA engine
- Allowed developing device/character driver against a generic DMA Linux API
- Required a patch to supply result with 'residue' populated for transfers that were less than the buffer size. A patch fixing this was applied 5 months ago.
Commit([0f7b82f11](#))
- Driver was unstable for long running DMA transfers, would have spontaneous failures, requiring reset of DMA Controller
- Several patches have been applied since 2017.4 release, which could address stability issues

▶ IIO Driver for Xilinx Linux Kernel/DMA Controller

- Analog Devices solution!
- Appealing option for common interface implementation
- Ultimately decided a simple clean solution tailored to our specific needs would be better

Custom solution for our application likely the quickest and most reliable solution to the problem

► Xilinx AXI4

- Implement
- Allowed to be implemented against
- Requires 'residual' less than was applied to Comm
- Driver transfer requiring
- Several 2017.4 issues

HOW STANDARDS PROLIFERATE:
(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)



Controller

Interface

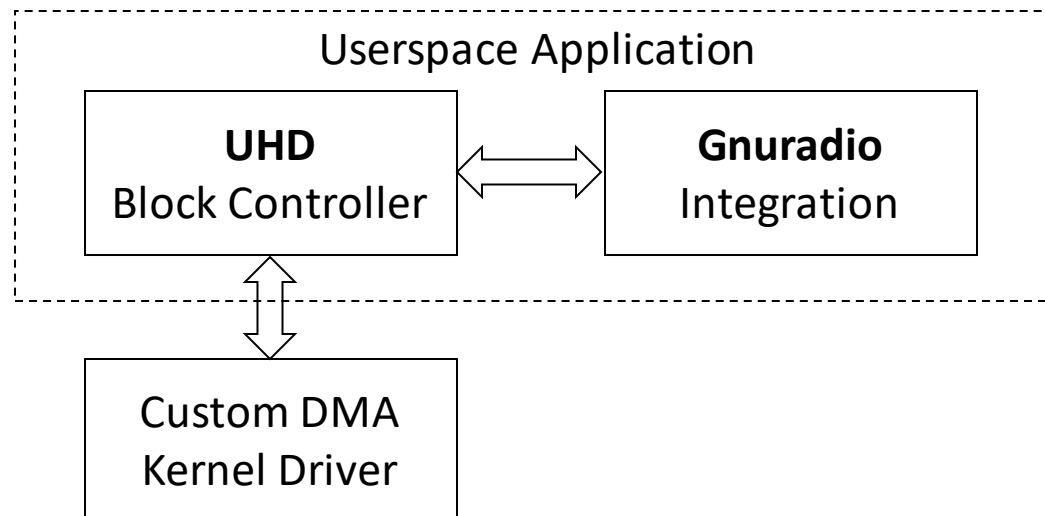
Custom solution
would be better

Custom solution for our application likely the quickest and most reliable solution to the problem

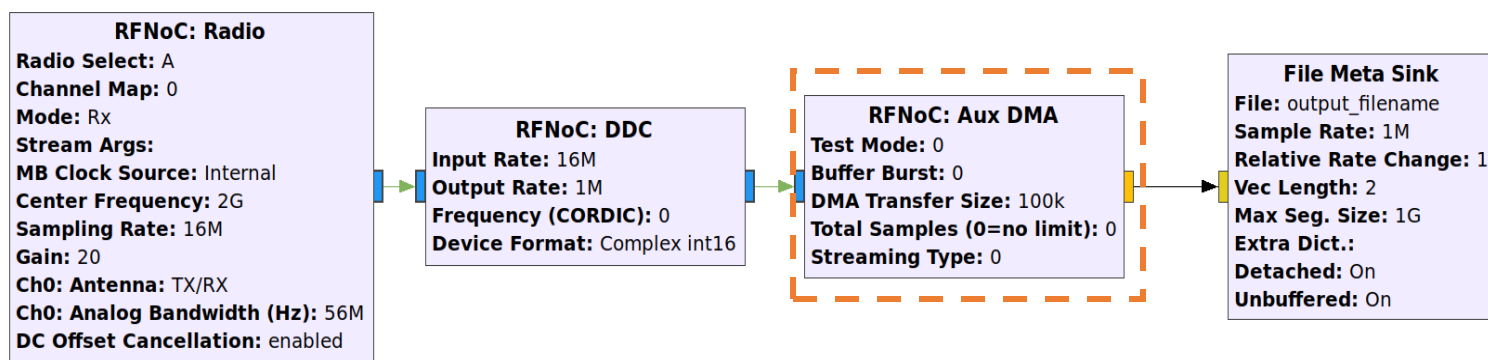
- ▶ Custom AUX DMA Linux Driver!
 - Implemented as a character device driver
 - Pre-allocate memory buffers `dma_alloc_coherent` when user space application calls open on file handle
 - *Kernel* parameters to specify how many buffers, and what size the buffers are.
 - When multiple buffers are allocated, will automatically swap free buffers for filled buffers in DMA interrupt, until all free buffers are filled, at which point DMA controller will be stopped.
- ▶ Multiple large buffers allows us to capture several seconds of high rate RF data uninterrupted into pre-allocated kernel buffers
- ▶ Allocate 160 MiB RAM for DMA
- ▶ Userspace applications then process/store after the burst has completed

Larger buffers decrease PS interaction load; Userspace can service more data!

- ▶ Userspace applications can retrieve buffers 2 ways: blocking **read()** or **mmap()**
 - **read()** will copy buffer to userspace memory, and automatically free kernel buffer. Not as efficient.
 - **mmap()** maps kernel buffer to userspace, and user application gets address of next filled buffer via ioctl call. Frees read buffers with ioctl call. ("zero-copy")
- ▶ Memory-mapped buffer handling performed in UHD block controller software

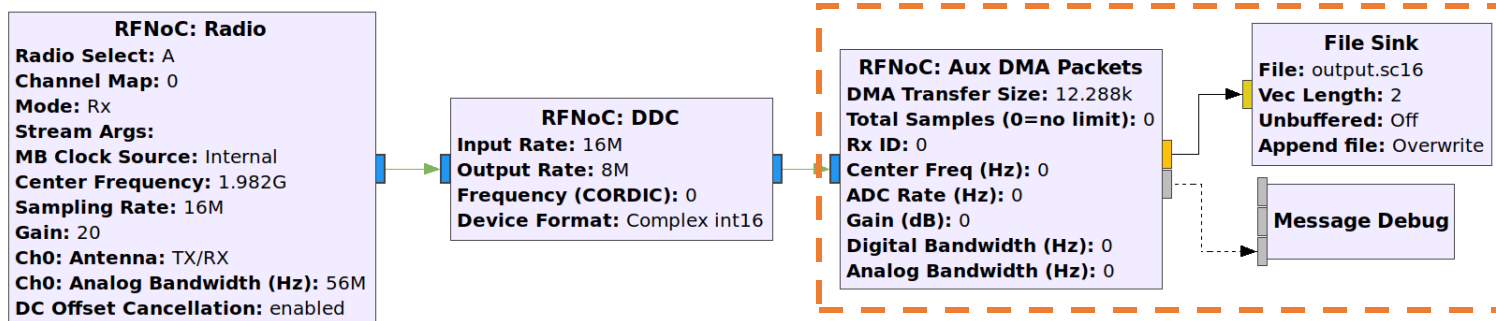



- ▶ Override `general_work` function of Gnuradio block
 - Request a new buffer from kernel driver
 - Fill the output buffer with the desired number of samples
 - Save samples greater than the requested `noutput_items` for the next iteration
- ▶ Optimized to transfer large buffers (128 KiB) straight into Gnuradio as fast as possible!
- ▶ Downsides: One channel only, zero flexibility, potentially large latency



Audience Inquiry: Better ways to transfer data to/from Gnuradio besides "general work"?

- ▶ Transfer **full CHDR packets** rather than streaming data
 - Skip "AXI Wrapper" in RFNoC Block
 - Pack many RFNoC transactions into one DMA transfer
 - Parse CHDR timestamp and metadata in Gnuradio
 - Convert output to PDUs or dump to disk
- ▶ Optimized for transferring many small packets or bursts
 - May also interleave multiple channels!



Get Creative! It's software!



Results & Tradeoffs



- ▶ (No PS Load) Max continuous receive rate without overflows: **40 MHz**
- ▶ (Writing to File) Max **continuous** receive rate without overflow: **~4 MHz**
 - 1-channel receive only
 - Collect continuously for 30 seconds, using Gnuradio's file meta sink
- ▶ (Writing to File) Max **burst**ed receive rate without overflow: **40 MHz**
 - First, write data into 160 MB allocated RAM
 - Second, write data from RAM into hard drive

"Continuous" rate limited by SD-card write speed and conflicting IO interrupts

"Burst" rate saves the full data buffer into RAM before dumping to disk

RFNoC Default DMA

- ▶ Small DMA buffers:
 - Lower **latency**
- ▶ Scatter Gather:
 - More **FPGA** control



HE360 DMA

- ▶ Large DMA buffers:
 - Higher **bandwidth**
- ▶ Direct Register Mode:
 - More **Processor** control
 - Simpler to implement

Consider **DMA implementation** vs **application goals**
when designing or selecting your transfer layer

Summary

- ▶ Heterogenous data transfer between processor and coprocessor (FPGA, GPU, etc) is a standard "problem" area
- ▶ Not all platforms/implementations are optimal for YOUR application
- ▶ Successful implementation requires a combination of custom FPGA, Linux kernel drivers, and userspace software
- ▶ We hope you feel equipped to address these issues in the future!

Full-rate heterogenous transfers into and out of Gnuradio can run on embedded Zynq-based SoC platforms



Thank You!

EJ Kreinar: ej@he360.com

Lorin Metzger: lorin@he360.com

Dan CaJacob: dan@he360.com

